
Detecting Corrupted Labels Without Training a Model to Predict

Zhaowei Zhu¹ Zihao Dong¹ Yang Liu¹

Abstract

Label noise in real-world datasets encodes wrong correlation patterns and impairs the generalization of deep neural networks (DNNs). It is critical to find efficient ways to detect corrupted patterns. Current methods primarily focus on designing robust training techniques to prevent DNNs from memorizing corrupted patterns. These approaches often require customized training processes and may overfit corrupted patterns, leading to a performance drop in detection. In this paper, from a more data-centric perspective, we propose a training-free solution to detect corrupted labels. Intuitively, “closer” instances are more likely to share the same clean label. Based on the neighborhood information, we propose two methods: the first one uses “local voting” via checking the noisy label consensuses of nearby features. The second one is a ranking-based approach that scores each instance and filters out a guaranteed number of instances that are likely to be corrupted. We theoretically analyze how the quality of features affects the local voting and provide guidelines for tuning neighborhood size. We also prove the worst-case error bound for the ranking-based method. Experiments with both synthetic and real-world label noise demonstrate our training-free solutions consistently and significantly improve most of the training-based baselines. Code is available at github.com/UCSC-REAL/SimiFeat.

1. Introduction

The generalization of deep neural networks (DNNs) depends on the quality and the quantity of the data. Nonetheless, real-world datasets often contain label noise that challenges the above assumption (Krizhevsky et al., 2012; Zhang et al., 2017; Agarwal et al., 2016; Wang et al., 2021a). Employing

human workers to clean annotations is one reliable way to improve the label quality, but it is too expensive and time-consuming for a large-scale dataset. One promising way to automatically clean up label errors is to first algorithmically detect possible label errors from a large-scale dataset (Cheng et al., 2021a; Northcutt et al., 2021a; Pruthi et al., 2020; Bahri et al., 2020), and then correct them using either algorithm or crowdsourcing (Northcutt et al., 2021b).

Almost all the algorithmic detection approaches focus on designing customized training processes to learn with noisy labels, where the idea is to train DNNs with noisy supervisions and then make decisions based on the output (Northcutt et al., 2021a) or gradients (Pruthi et al., 2020) of the last logit layer of the trained model. The high-level intuition of these methods is the memorization effects (Han et al., 2020), i.e., instances with label errors, a.k.a., corrupted instances, tend to be harder to be learned by DNNs than clean instances (Xia et al., 2021; Liu et al., 2020; Bai & Liu, 2021). By setting appropriate hyperparameters to utilize the memorization effect, corrupted instances could be identified.

Limitations of the learning-centric methods The above methods suffer from *two major limitations*: 1) the customized training processes are task-specific and may require fine-tuning hyperparameters for different datasets/noise; 2) as long as the model is trained with noisy supervisions, the memorization of corrupted instances exists. The model will “subjectively” and wrongly treat the memorized/overfitted corrupted instances as clean. For example, some low-frequency/rare clean instances may be harder to memorize than high-frequency/common corrupted instances. Memorizing these corrupted instances lead to unexpected and disparate impacts (Liu, 2021).

Existing solutions to avoid memorizing/overfitting are to employ some regularizers (Cheng et al., 2021a) or use early-stopping (Bai et al., 2021; Li et al., 2020b). However, their performance depends on hyperparameter settings. One promising way to avoid memorizing/overfitting is to drop the dependency on the noisy supervision, which motivates us to design a *training-free* method to find label errors. Intuitively, we can carefully use the information from nearby features to infer whether one instance is corrupted or not. The comparison between our data-centric and

¹Department of Computer Science and Engineering, University of California, Santa Cruz, CA, USA. Correspondence to: Yang Liu <yangliu@ucsc.edu>.




Figure 1. The existing learning-centric pipeline vs. our proposed data-centric pipeline. The inputs are features and the corresponding noisy labels, and the outputs are a set of corrupted labels. **Blue:** The learning-centric solution. **Orange:** The data-centric solution.

existing learning-centric solutions is illustrated in Figure 1.

Our training-free method enables more possibilities beyond a better detection result. For example, the concerns about the required assumptions and hyperparameter tuning in those training-based methods will now be released due to our training-free property. The complexity will also be much lower, again due to removing the possibly involved training processes. This light detection solution also has the potential to serve as a pre-processing module to prepare data for other sophisticated tasks (e.g., semi-supervised learning (Xie et al., 2019; Berthelot et al., 2019)).

Our main contributions are:

- *New perspective:* Different from current methods that train customized models on noisy datasets, we propose a training-free and data-centric solution to efficiently detect corrupted labels, which provides a new and complementary perspective to the traditional learning-centric solution. We demonstrate the effectiveness of this simple idea and open the possibility for follow-up works.
- *Efficient algorithms:* Based on the neighborhood information, we propose two methods: a voting-based local detection method that only requires checking the noisy label consensuses of nearby features, and a ranking-based global detection method that scores each instance by its likelihood of being clean and filters out a guaranteed percentage of instances with low scores as corrupted ones.
- *Theoretical analyses:* We theoretically analyze how the quality of features (but possibly imperfect in practice) affects the local voting and provide guidelines for tuning neighborhood size. We also prove the worst-case error bound for the ranking-based method.
- *Numerical findings:* Our numerical experiments show three important messages: in corrupted label detection, i) training with noisy supervisions may not be necessary; ii) feature extraction layers tend to be more useful than the logit layers; iii) features extracted from other tasks or domains are helpful.

1.1. Related Works

Learning with noisy labels There are many other works that can detect corrupted instances (a.k.a. sample selection) in the literature, e.g., (Han et al., 2018; Yu et al., 2019; Yao et al., 2020; Wei et al., 2020; Jiang et al., 2020; Zhang et al., 2021; Huang et al., 2019), and its combination with semi-supervised learning (Wang et al., 2020; Li et al., 2020a; Cheng et al., 2021a). Another line of works focus on designing robust loss functions to mitigate the effect of label noise, such as numerical methods (Ghosh et al., 2017; Zhang & Sabuncu, 2018; Gong et al., 2018; Amid et al., 2019; Wang et al., 2019; Shu et al., 2020; Wang et al., 2022a) and statistical methods (Natarajan et al., 2013; Liu & Tao, 2015; Patrini et al., 2017; Liu & Guo, 2020; Xia et al., 2019; Zhu et al., 2021a; Jiang et al., 2022; Feng et al., 2021; Wei et al., 2021b; 2022c). They all require training DNNs with noisy supervisions and would suffer from the limitations of the learning-centric methods.

k -NN for noisy labels The k -NN technique often plays important roles in building auxiliary methods to improve deep learning (Jiang et al., 2018). Recently, it has been extended to filtering out corrupted instances when learning with noisy labels (Gao et al., 2016; Reeve & Kabán, 2019; Kong et al., 2020; Bahri et al., 2020). However, these methods focus on learning-centric solutions and cannot avoid memorizing noisy labels.

Label aggregation Our work is also relevant to the literature of crowdsourcing that focuses on label aggregation (to clean the labels) (Liu et al., 2012; Karger et al., 2011; 2013; Liu & Liu, 2015; Zhang et al., 2014; Wei et al., 2022e). Most of these works can access multiple reports (labels) for the same input feature, while our real-world datasets usually have only one noisy label for each feature.

2. Preliminaries

Instances Traditional supervised classification tasks build on a clean dataset $D := \{(x_n, y_n)\}_{n \in [N]}$, where $[N] := \{1, 2, \dots, N\}$. Each *clean instance* (x_n, y_n) includes *feature* x_n and *clean label* y_n , which is drawn according to random variables $(X, Y) \sim \mathcal{D}$. In many practical cases, the clean labels may be unavailable and the learner could only observe a noisy dataset denoted by $\tilde{D} := \{(x_n, \tilde{y}_n)\}_{n \in [N]}$, where (x_n, \tilde{y}_n) is a *noisy instance* and the *noisy label* \tilde{y}_n may or may not be identical to y_n . We call \tilde{y}_n is *corrupted* if $\tilde{y}_n \neq y_n$ and clean otherwise. The instance (x_n, \tilde{y}_n) is a *corrupted instance* if \tilde{y}_n is corrupted. The noisy data distribution corresponds to \tilde{D} is $(X, \tilde{Y}) \sim \tilde{\mathcal{D}}$. We focus on the closed-set label noise that Y and \tilde{Y} are assumed to be in the same label space, e.g., $Y, \tilde{Y} \in [K]$. Explorations on open-set data (Xia et al., 2020a; Wei et al., 2021a; Luo et al., 2021) are deferred to future works.

Clusterability In this paper, we focus on a setting where the distances between two features should be comparable or clusterable (Zhu et al., 2021b), i.e., nearby features should belong to the same true class with a high probability (Gao et al., 2016), which could be formally defined as:

Definition 2.1 ((k, δ_k) label clusterability). A dataset D satisfies (k, δ_k) label clusterability if: $\forall n \in [N]$, the feature x_n and its k -Nearest-Neighbors (k -NN) x_{n_1}, \dots, x_{n_k} belong to the same true class with probability at least $1 - \delta_k$.

Note δ_k captures two types of randomnesses: one comes from a probabilistic Y given X , i.e., $\exists i, x, \mathbb{P}(Y = i | X = x) \notin \{0, 1\}$; the other depends on the quality of features and the value of k , which will be further illustrated in Figure 3. The $(k, 0)$ label clusterability is also known as k -NN label clusterability (Zhu et al., 2021b).

Corrupted label detection Our paper aims to improve the performance of the corrupted label detection (a.k.a. finding label errors) which is measured by the F_1 -score of the *detected corrupted instances*, which is the harmonic mean of the precision and recall, i.e.

$$F_1 = 2 / (\text{Precision}^{-1} + \text{Recall}^{-1}).$$

Let $\mathbb{1}(\cdot)$ be the indicator function that takes value 1 when the specified condition is satisfied and 0 otherwise. Let $v_n = 1$ indicate that \tilde{y}_n is detected as a corrupted label, and $v_n = 0$ if \tilde{y}_n is detected to be clean. Then the precision and recall can be calculated as

$$\text{Precision} = \frac{\sum_{n \in [N]} \mathbb{1}(v_n = 1, \tilde{y}_n \neq y_n)}{\sum_{n \in [N]} \mathbb{1}(v_n = 1)},$$

$$\text{Recall} = \frac{\sum_{n \in [N]} \mathbb{1}(v_n = 1, \tilde{y}_n \neq y_n)}{\sum_{n \in [N]} \mathbb{1}(\tilde{y}_n \neq y_n)}.$$

Note the F_1 score on corrupted instances is sensitive to the case when the noise rate is mild to low, which is typically the case in practice. For example, if 20% of the data is corrupted but the algorithm reports no label errors, the returned F_1 score on corrupted instances is 0 while the one on clean instances is $2 / (0.8^{-1} + 1) \approx 0.89$.

3. Corrupted Label Detection Using Similar Features

Different from most methods that detect corrupted instances based on the logit layer or model predictions (Northcutt et al., 2021a; Cheng et al., 2021a; Pruthi et al., 2020; Bahri et al., 2020), we focus on a more data-centric solution that operates on features. Particularly we are interested in the possibility of detecting corrupted labels in a training-free way. In this section, we will first introduce intuitions, and then provide two efficient algorithms to detect corrupted labels with similar features.




Figure 2. Detect corrupted labels with similar features. Orange circle: instance with noisy label 1. Blue square: instance with noisy label 2. Green dashed circle: A k -NN example.

3.1. Intuitions

The learning-centric detection methods often detect corrupted instances by comparing model predictions with noisy labels (Cheng et al., 2021a; Northcutt et al., 2021a) as illustrated in Figure 1. However, for the data-centric method, the feature x_n cannot be directly compared with the noisy label \tilde{y}_n since x_n is not directly categorical without a model, i.e., the connection between a single x_n and \tilde{y}_n is weak. Thus our first step should be establishing an auxiliary categorical information using only features.

As illustrated in Figure 2, the high-level intuition is to check label consensus of nearby features. With $(k, 0)$ label clusterability as in Definition 2.1, we know the true labels of x_n and its k -NN x_{n_1}, \dots, x_{n_k} should be the same. If we further assume the label noise is group-dependent (Wang et al., 2021b), i.e., each x_n and its k -NN can be viewed as a local group and share the same noise transition matrix (Liu, 2022): $\mathbb{P}(\tilde{Y} = \tilde{y}_n | X = x_n, Y = y_n) = \mathbb{P}(\tilde{Y} = \tilde{y}_n | X = x_{n'}, Y = y_n), \forall x_{n'} \in \{x_{n_1}, \dots, x_{n_k}\}$, we can first treat their noisy labels as $k + 1$ independent observations of $\mathbb{P}(\tilde{Y} = \tilde{y}_n | X = x_n, Y = y_n)$, then estimate the probability by counting the (weighted) frequency of each class in the k -NN label estimator, and get k -NN labels \hat{y}_n . We use the **bold y** to indicate a vector, which can be seen as either an one-hot encoding of a hard label or a soft label (Zhu et al., 2022a). The i -th element $\hat{y}_n[i]$ can be interpreted as the estimated probability of predicting class- i .

Note the k -NN technique has been implemented by Bahri et al. (2020) as a filter to remove corrupted instances. However, this approach focuses on calculating distances on the logit layer, which inevitably requires a task-specific training process and may suffer from the limitations mentioned in Section 1. Besides, using appropriate features may be better than model logits/predictions when the dataset is noisy. See discussions below.

Features could be better than model predictions During supervised training, memorizing noisy labels makes the model generalize poorly (Han et al., 2020), while using only features may effectively avoid this issue (Li et al., 2021a). For those pre-extracted features, e.g., tabular data in UCI datasets (Dua & Graff, 2017), the input features are already comparable and directly applying data-centric methods on these features avoids memorizing noisy labels. For more challenging tasks such as image or text classifications, we can also borrow some pre-trained models to pre-process the raw feature to improve the clusterability of features, such as BERT (Devlin et al., 2019) for language tasks, CLIP (Radford et al., 2021) for vision-language tasks, or some feature extractors from unsupervised learning (Ji et al., 2019) and self-supervised learning (Jaiswal et al., 2021; Liu et al., 2021a; He et al., 2020; Chen et al., 2020; Cheng et al., 2021b), which are not affected by noisy labels.

3.2. Voting-Based Local Detection

Inspired by the idea implemented in model decisions, i.e., selecting the most likely class as the true class, we can simply “predict” the index that corresponds to the largest element in \hat{y}_n with random tie-breaking, i.e., $y_n^{\text{vote}} = \arg \max_{i \in [K]} \hat{y}_n[i]$. To further detect whether \hat{y}_n is corrupted or not, we only need to check $v_n := \mathbb{1}(y_n^{\text{vote}} \neq \tilde{y}_n)$. Recall $v_n = 1$ indicates a corrupted label. This voting method relies only on the local information within each k -NN label \hat{y}_n , which may not be robust with low-quality features. Intuitively, when the gap between the true class probability and the wrong class probability is small, the majority vote will be likely to make mistakes due to sampling errors in \hat{y}_n . Thus only using local information within each \hat{y}_n may not be sufficient. It is important to leverage more information such as some global statistics, which will be discussed later.

3.3. Ranking-Based Global Detection

The score function can be designed to detect corrupted instances (Northcutt et al., 2021a; Cheng et al., 2021a; Pruthi et al., 2020; Bahri et al., 2020), hard-to-learn instances (Liu et al., 2021b), out-of-distribution instances (Wei et al., 2022b), and suspicious instances that may cause model unfairness (Wang et al., 2022b). However, it is not clear how to do so without training a task-specific model. From a global perspective, if the likelihood for each instance being clean could be evaluated by some scoring functions, we can sort the scores in an increasing order and filter out the low-score instances as corrupted ones. Based on this intuition, there are two critical components: the *scoring function* and the *threshold* to differentiate the low-score part (corrupted) and the high-score part (clean).

Scoring function A good scoring function should be able to

give clean instances higher scores than corrupted instances. We adopt cosine similarity defined as:

$$\text{Score}(\hat{y}_n, j) = \frac{\hat{y}_n^\top e_j}{\|\hat{y}_n\|_2 \|e_j\|_2},$$

where e_j is the one-hot encoding of label j . To evaluate whether the soft label \hat{y}_n informs us a clean instance or not, we compare $\text{Score}(\hat{y}_n, \tilde{y}_n)$ with other instances that have the same noisy label. This scoring function captures more information than majority votes, which is summarized as follows.

Property 3.1 (Relative score). Within the same instance, the score of the majority class is higher than the others, i.e., $\forall j \neq y_n^{\text{vote}}, j \in [K], \forall n \in [N] : \text{Score}(\hat{y}_n, y_n^{\text{vote}}) > \text{Score}(\hat{y}_n, j)$.

Property 3.2 (Absolute score). $\text{Score}(\hat{y}_n, j)$ is jointly determined by both $\hat{y}_n[j]$ and $\hat{y}_n[j']$, $\forall j' \neq j$.

The first property guarantees that the corrupted labels would have lower scores than clean labels for the same instance when the vote is correct. However, although solely relying on Property 3.1 may work well in the voting-based method which makes decisions individually for each instance, it is not sufficient to be trustworthy in the ranking-based global detection. Empirically we observe that if we choose a score function that Property 3.2 does not hold, e.g., treating k -NN soft labels as model predictions and check the cross-entropy loss, it does not always return satisfying results in our experiments. The main reason is that, across different instances, the non-majority classes of some instances may have higher absolute scores than the majority classes of the other instances, which is especially true for general instance-dependent label noise with heterogeneous noise rates (Cheng et al., 2021a). Property 3.2 helps make it less likely to happen. Consider an example as follows.

Example Suppose $\hat{y}_{n_1} = \hat{y}_{n_2} = [0.6, 0.4, 0.0]^\top$, $\hat{y}_{n_3} = [0.34, 0.33, 0.33]^\top$, $y_{n_1} = y_{n_2} = y_{n_3} = 1$, $\tilde{y}_{n_1} = \tilde{y}_{n_3} = 1$, $\tilde{y}_{n_2} = 2$. We can use the majority vote to get perfect detection in this case, i.e., $y_{n_1}^{\text{vote}} = y_{n_2}^{\text{vote}} = y_{n_3}^{\text{vote}} = 1 = y_{n_1}$, since the first class of each instance has the largest value. However, if we directly use a single value in soft label \hat{y}_n to score them, e.g., $\text{Score}'(\hat{y}_n, j) = \hat{y}_n[j]$, we will have $\hat{y}_{n_1}[\tilde{y}_{n_1}] = 0.6 > \hat{y}_{n_2}[\tilde{y}_{n_2}] = 0.4 > \hat{y}_{n_3}[\tilde{y}_{n_3}] = 0.34$, where the ranking is $n_3 \prec n_2 \prec n_1$. Ideally, we know instance n_2 is corrupted and the true ranking should be $n_2 \prec n_3 \prec n_1$ or $n_2 \prec n_1 \prec n_3$. To mitigate this problem, we choose the cosine similarity as our scoring function. The three instances could be scored as 0.83, 0.55, 0.59, corresponding to an ideal ranking $n_2 \prec n_3 \prec n_1$. We formally introduce the detailed ranking approach as follows.

Ranking Suppose we have a group of instances with the same noisy class j , i.e. $\{(x_n, \tilde{y}_n)\}_{n \in \mathcal{N}_j}$, where $\mathcal{N}_j := \{n | \tilde{y}_n = j\}$ are the set of indices that correspond to noisy

class j . Let N_j be the number of indices in \mathcal{N}_j (counted from noisy labels). Intuitively, we can first sort all instances in \mathcal{N}_j in an increasing order by `argsort` and obtain the original indices for the sorted scores as:

$$\mathcal{I} = \text{argsort}\{\text{Score}(\hat{y}_n, j)\}_{n \in \mathcal{N}_j},$$

where the low-score head is supposed to consist of corrupted instances (Northcutt et al., 2021a). Then we can simply select the first \tilde{N}_j instances with low scores as corrupted instances:

$$v_n = \mathbb{1}(\text{Loc}(n, \mathcal{I}) \leq \tilde{N}_j),$$

where $\text{Loc}(n, \mathcal{I})$ returns the index of n in \mathcal{I} . Instead of manually tuning \tilde{N}_j (Han et al., 2018), we discuss how to determine it algorithmically.

Threshold The number of corrupted instances in \mathcal{N}_j is approximately $\mathbb{P}(Y \neq j | \tilde{Y} = j) \cdot N_j$ when N_j is sufficiently large. Therefore if all the corrupted instances have lower scores than any clean instance, we can set $\tilde{N}_j = \mathbb{P}(Y \neq j | \tilde{Y} = j) \cdot N_j$ to obtain the ideal division. Note N_j can be obtained by directly counting the number of instances with noisy label j . To calculate the probability

$$\mathbb{P}(Y \neq j | \tilde{Y} = j) = 1 - \mathbb{P}(Y = j | \tilde{Y} = j),$$

we borrow the results from the HOC estimator (Zhu et al., 2021b; 2022b), where the noise transition probability $\mathbb{P}(\tilde{Y} = j | Y = j)$ and the marginal distribution of clean label $\mathbb{P}(Y = j)$ can be estimated with only features and the corresponding noisy labels. Then we can calculate our needed probability by Bayes' rule

$$\mathbb{P}(Y = j | \tilde{Y} = j) = \mathbb{P}(\tilde{Y} = j | Y = j) \cdot \mathbb{P}(Y = j) / \mathbb{P}(\tilde{Y} = j),$$

where $\mathbb{P}(\tilde{Y} = j)$ can be estimated by counting the frequency of noisy label j in \tilde{D} . Technically other methods exist in the literature to estimate $\mathbb{P}(\tilde{Y} | Y)$ (Liu & Tao, 2015; Patrini et al., 2017; Northcutt et al., 2021a; Li et al., 2021b). But they often require training a model to fit the data distribution, which conflict with our goal of a training-free solution; instead, HOC fits us perfectly.

3.4. Algorithm: SimiFeat

Algorithm 1 summarizes our solution. The main computation complexity is pre-processing features with extractor $g(\cdot)$, which is less than the cost of evaluating the model compared with the training-based methods. Thus SimiFeat can filter out corrupted instances efficiently. In Algorithm 1, we run either voting-based local detection as Lines 1, 1, or ranking-based global detection as Lines 1, 1. The detection is run multiple times with random standard data augmentations to reduce the variance of estimation. The majority of results from different epochs is adopted as the final detection output as Line 1, i.e., flag as corrupted if $v_n = 1$ in more than half of the epochs.

4. How Does Feature Quality Affect Our Solution?

In this section, we will first show how the quality of features¹ affects the selection of the hyperparameter k , then analyze the error upper bound for the ranking-based method.

4.1. How Does Feature Quality Affect the Choice of k ?

Recall k is used as illustrated in Figure 2. On one hand, the k -NN label estimator will be more accurate if there is stronger clusterability that more neighbor features belong to the same true class (Liu & Liu, 2015; Zhu et al., 2021b), which helps improve the performance of later algorithms. On the other hand, with good but imperfect features, stronger clusterability with a larger k is less likely to satisfy, thus the violation probability δ_k increases with k for a given extractor $g(\cdot)$. We take the voting-based method as an example and analyze this tradeoff. For a clean presentation, we focus on a binary classification with instance-dependent label noise where $\mathbb{P}(Y = 1) = p$, $\mathbb{P}(\tilde{Y} = 2 | X, Y = 1) = e_1(X)$, $\mathbb{P}(\tilde{Y} = 1 | X, Y = 2) = e_2(X)$. Suppose the instance-dependent noise rate is upper-bounded by e , i.e., $e_1(X) \leq e$, $e_2(X) \leq e$. With δ_k as in Definition 2.1, we calculate the lower bound of the probability that the vote is correct in Proposition 4.1.

Proposition 4.1. *The lower bound for the probability of getting true detection with majority vote is*

$$\mathbb{P}(\text{Vote is correct} | k) \geq (1 - \delta_k) \cdot I_{1-e}(k + 1 - k', k' + 1),$$

where $k' = \lceil (k + 1)/2 \rceil - 1$, $I_{1-e}(k + 1 - k', k' + 1)$ is the regularized incomplete beta function defined as $I_{1-e}(k + 1 - k', k' + 1) = (k + 1 - k') \binom{k+1}{k'} \int_0^{1-e} t^{k-k'} (1-t)^{k'} dt$.

Proposition 4.1 shows the tradeoff between a reliable k -NN label and an accurate vote. When k is increasing, **Term-1** $(1 - \delta_k)$ (quality of features) decreases but **Term-2** $I_{1-e_1}(k + 1 - k', k' + 1)$ (result of pure majority vote) increases. With Proposition 4.1, we are ready to answer the question: *when do we need more labels?* See Remark 4.2.

Remark 4.2. Consider the lower bounds with k_1 and k_2 ($k_1 < k_2$). Supposing the first lower bound is lower than the second lower bound, based on Proposition 4.1, we roughly study the trend with an increasing k by comparing two bounds and get

$$\frac{1 - \delta_{k_1}}{1 - \delta_{k_2}} < \frac{I_{1-e}(k_2 + 1 - k'_2, k'_2 + 1)}{I_{1-e}(k_1 + 1 - k'_1, k'_1 + 1)}.$$

For example, when $k_1 = 5$, $k_2 = 20$, $e = 0.4$, we can calculate the incomplete beta function and $\frac{1-\delta_5}{1-\delta_{20}} < 1.52$. Supposing $\delta_5 = 0.2$, we have $\delta_{20} < 0.47$. This indicates

¹Note the voting-based method achieves an F_1 -score of 1 when $(k, 0)$ -NN label clusterability, $k \rightarrow +\infty$, holds.

Algorithm 1 Detection with **Similar Features** (The SimiFeat Detector)

```

1: Input: Number of epochs:  $M$ .  $k$ -NN parameter:  $k$ . Noisy dataset:  $\tilde{D} = \{(x_n, \tilde{y}_n)\}_{n \in [N]}$ . Feature extractor:  $g(\cdot)$ .
   Method: Vote or Rank. Epoch counter  $m = 0$ .
2: repeat
3:    $x'_n \leftarrow \text{RandPreProcess}(x_n), \forall n;$  # Initialize & Standard data augmentations
4:    $x_n \leftarrow g(x'_n), \forall n;$  # For tasks with rarely clusterable features, extract features with  $g(\cdot)$ 
5:    $\hat{y}_n \leftarrow \text{kNNLabel}(\{x_n\}_{n \in [N]}, k)$  # Get soft labels. One can weight instances by the similarity to the center instance.
6:   if Vote then
7:      $y_n^{\text{vote}} \leftarrow \arg \max_{i \in [K]} \hat{y}_n[i];$  # Apply local majority vote
8:      $v_n \leftarrow \mathbb{1}(y_n^{\text{vote}} \neq \tilde{y}_n), \forall n \in [N];$  # Treat as corrupted if majority votes disagree with noisy labels
9:   else
10:     $\mathbb{P}(Y), \mathbb{P}(\tilde{Y}|Y) \leftarrow \text{HOC}(\{(x_n, \tilde{y}_n)\}_{n \in [N]});$  # Estimate clean priors  $\mathbb{P}(Y)$  and noise transitions  $\mathbb{P}(\tilde{Y}|Y)$  by HOC
11:     $\mathbb{P}(Y|\tilde{Y}) = \mathbb{P}(\tilde{Y}|Y) \cdot \mathbb{P}(Y) / \mathbb{P}(\tilde{Y});$  # Estimate thresholds by Bayes' rule
12:    for  $j$  in  $[K]$  do
13:       $\mathcal{N}_j := \{n | \tilde{y}_n = j\};$  # Detect corrupted labels in each set  $\mathcal{N}_j$ 
14:       $\mathcal{I} \leftarrow \text{argsort}\{\text{Score}(\hat{y}_n, j)\}_{n \in \mathcal{N}_j};$  #  $\mathcal{I}$  records the raw index of each sorted value
15:       $v_n \leftarrow \mathbb{1}(\text{Loc}(n, \mathcal{I}) \leq \lfloor (1 - \mathbb{P}(Y = j | \tilde{Y} = j)) \cdot N_j \rfloor);$  # Select low-score (head) instances as corrupted ones
16:    end for
17:  end if
18:   $\mathcal{V}_m = \{v_n\}_{n \in [N]};$  # Record detection results in the  $m$ -th epoch
19: until  $M$  times
20:  $\mathcal{V} = \text{Vote}(\mathcal{V}_m, \forall m \in [M]);$  # Do majority vote based on results from  $M$  epochs
21: Output:  $[N] \setminus \mathcal{V}$ .

```

increasing k from 5 to 20 would not improve the lower bound with features satisfying $\delta_{20} > 0.47$. This observation helps us set k with practical and imperfect features. We set $k = 10$ in all of our experiments.

Remark 4.2 indicates that: *with practical (imperfect) features, a small k may achieve the best (highest) probability lower bound*. To further consolidate this claim, we numerically calculate δ_k with different quality of features on CIFAR-10 and the corresponding probability lower bound in Figure 3. We find most of the probability lower bounds first increase then decrease except for the “perfect” feature which is extracted by the extractor trained using ground-truth labels. Note this feature extractor has memorized all clean instances so that $\delta_k \rightarrow 0$ since $k \ll 5000$ (the number of instances in the same label class).

4.2. How Does Feature Quality Affect F_1 -Score?

We next prove the probability bound for the performance of the ranking-based method. Consider a K -class classification problem with informative instance-dependent label noise (Cheng et al., 2021a). Denote random variable S by the score of each instance being clean. A higher score S indicates the instance is more likely to be clean. Denote the score of a true/false instance by $S_{n,j}^{\text{true}} := \text{Score}(\hat{y}_n, j)$ when $\tilde{y}_n = y_n = j$ and $S_{n',j}^{\text{false}} := \text{Score}(\hat{y}_{n'}, j)$ when $\tilde{y}_{n'} = j, y_{n'} \neq j$. Both are scalars. Then for instances in \mathcal{N}_j , we have two set of random variables $\mathbb{S}_j^{\text{true}} := \{S_{n,j}^{\text{true}} | n \in$

$\mathcal{N}_j, \tilde{y}_n = y_n = j\}$ and $\mathbb{S}_j^{\text{false}} := \{S_{n',j}^{\text{false}} | n' \in \mathcal{N}_j, \tilde{y}_{n'} = j, y_{n'} \neq j\}$. Recall $\mathcal{N}_j := \{n | \tilde{y}_n = j\}$ are the set of indices that correspond to noisy class j . Intuitively, the score $S_{n,j}^{\text{true}}$ should be greater than $S_{n',j}^{\text{false}}$. Suppose their means, which depend on noise rates, are bounded, i.e.,

$$\mathbb{E}[S_{n,j}^{\text{true}}] \geq \mu_j^{\text{true}}, \quad \mathbb{E}[S_{n',j}^{\text{false}}] \leq \mu_j^{\text{false}}$$

for all feasible n, n' . Assume there exists a feasible v such that both S_j^{true} and S_j^{false} follow sub-Gaussian distributions with variance proxy $\frac{\Delta^2}{2v}$ (Buldygin & Kozachenko, 1980; Zhu et al., 2021c) such that:

$$\mathbb{P}(\mu_j^{\text{true}} - S_{n,j}^{\text{true}} \geq t) \leq e^{-\frac{vt^2}{\Delta^2}}, \quad \mathbb{P}(S_{n',j}^{\text{false}} - \mu_j^{\text{false}} \geq t) \leq e^{-\frac{vt^2}{\Delta^2}},$$

and the probability density satisfies $\mathbb{P}(S_j^{\text{true}} = \mu_j^{\text{true}}) = \mathbb{P}(S_j^{\text{false}} = \mu_j^{\text{false}}) = 1/\Delta$, where $1/\Delta$ is the “height” of both distributions, v is the decay rate of tails. Let N_j^- (N_j^+) be the number of indices in $\mathbb{S}_j^{\text{false}}$ ($\mathbb{S}_j^{\text{true}}$). Theorem 4.3 summarizes the performance bound of the ranking-based method. See Appendix for the proof.

Theorem 4.3. *With probability at least p , when the threshold for the ranking-based method is set to*

$$(1 - \mathbb{P}(Y = j | \tilde{Y} = j)) \cdot N_j$$

as Line 15, the F_1 -score of detecting corrupted instances in \mathcal{N}_j by ranking is at least $1 - \frac{e^{-v} \max(N_j^-, N_j^+) + \alpha}{N_j^-}$, where $p =$




Figure 3. The trends of δ_k and probability lower bounds on CIFAR-10 (Krizhevsky et al., 2009) when raw features are extracted with different $g(\cdot)$. The outputs of the last convolution layer are adopted. R18/34/50: ResNet18/34/50. Img: Pre-trained on ImageNet (Deng et al., 2009). C10/100-Clean: Pre-trained on clean CIFAR-10/100. C10/100-SSL: Pre-trained on CIFAR-10/100 without labels by SimCLR (Chen et al., 2020). ViT-B/32-CLIP: CLIP Pre-trained vision transformer (Radford et al., 2021).

$\int_{-1}^{\mu_j^{\text{true}} - \mu_j^{\text{false}} - \Delta} f(t) dt$, $f(t)$ is the probability density function of the difference of two independent beta-distributed random variables $\beta_1 - \beta_2$, where $\beta_1 \sim \text{Beta}(N^-, 1)$, $\beta_2 \sim \text{Beta}(\alpha + 1, N^+ - \alpha)$.

Theorem 4.1 shows the detection performance depends on:

- the concentration of $S_{n,j}^{\text{true}}$ and $S_{n',j}^{\text{false}}$: variance proxy $\frac{\Delta^2}{2v}$;
- the distance between $S_{n,j}^{\text{true}}$ and $S_{n',j}^{\text{false}}$: $\mu_j^{\text{true}} - \mu_j^{\text{false}}$.

Intuitively, with proper scoring function and high-quality features, we have small variance proxy (small Δ and large v) and F_1 -score approximates to 1.

5. Empirical Results

We present experimental evidence in this section. The performance is measured by the F_1 -score of the detected corrupted labels as defined in Section 2. Note there is no training procedure in our method. The only hyperparameters in our methods are the number of epochs M and the k -NN parameter k . Intuitively, a larger M returns a collective result from more times of detection, which should be more accurate. But a larger M takes more time. We set $M = 21$ (an odd number for better tie-breaking) for an efficient solution. The hyperparameter k cannot be set too large as demonstrated in Figure 3. From Figure 3, we notice that the lower bound (RHS figure) is relatively high when $k = 10$ for all settings. Therefore, in CIFAR (Krizhevsky et al., 2009) experiments, rather than fine-tune M and k for different settings, we fix $M = 21$ and $k = 10$. We also test on Clothing1M (Xiao et al., 2015). Detailed experiment settings on Clothing1M are in Appendix C.

Synthetic label noise We experiment with three popular synthetic label noise models: the *symmetric* label noise, the *asymmetric* label noise, and the *instance-dependent* label noise. Denote the ratio of instances with corrupted labels in the whole dataset by η . Both the symmetric and the asymmetric noise models follow the class-dependent assumption

(Liu & Tao, 2015), i.e., the label noise only depends only on the clean class: $\mathbb{P}(\tilde{Y}|X, Y) = \mathbb{P}(\tilde{Y}|Y)$. Specially, the symmetric noise is generated by uniform flipping, i.e., randomly flipping a true label to the other possible classes w.p. η (Cheng et al., 2021a). The asymmetric noise is generated by pair-wise flipping, i.e., randomly flipping true label i to the next class $(i \bmod K) + 1$. Denote by d the dimension of features. The instance-dependent label noise is synthesized by randomly generating a $d \times K$ projection matrix w_i for each class i and project each incoming feature with true class y_n onto each column of w_{y_n} (Xia et al., 2020b). Instance n is more likely to be flipped to class j if the projection value of x_n on the j -th column of w_{y_n} is high. See Appendix B in (Xia et al., 2020b) and Appendix D.1 in (Zhu et al., 2021b) for more details. We use symmetric noise with $\eta = 0.6$ (*Symm. 0.6*), asymmetric noise with $\eta = 0.3$ (*Asym. 0.3*), and instance-dependent noise with $\eta = 0.4$ (*Inst. 0.4*) in experiments.

Real-world label noise The *real-world* label noise comes from human annotations or weakly labeled web data. We use the 50,000 noisy training labels ($\eta \approx 0.16$) for CIFAR-10 collected by (Zhu et al., 2021b), and 50,000 noisy training labels ($\eta \approx 0.40$) for CIFAR-100 collected by (Wei et al., 2022d). Both sets of noisy labels are crowd-sourced from Amazon Mechanical Turk. For Clothing1M (Xiao et al., 2015), we could not calculate the F_1 -scores due to the lack of ground-truth labels. We firstly perform noise detection on 1 million noisy training instances then train only with the selected clean data to check the effectiveness.

5.1. Fitting Noisy Distributions May Not Be Necessary

Our first experiment aims to show that fitting the noisy data distribution may not be necessary in detecting corrupted labels. To this end, we compare our methods, i.e., voting-based local detection (SimiFeat-V) and ranking-based global detection (SimiFeat-R), with three learning-centric noise detection works: CORES (Cheng et al., 2021a),

Table 1. Comparisons of F_1 -scores (%). CORES, CL, TracIn: Train with noisy supervisions. SimiFeat-V and SimiFeat-R: Get $g(\cdot)$ without any supervision. Top 2 are **bold**.

METHOD	CIFAR10				CIFAR100			
	Human	Symm. 0.6	Asym. 0.3	Inst. 0.4	Human	Symm. 0.6	Asym. 0.3	Inst. 0.4
CORES	65.00	92.94	7.68	87.43	3.52	92.34	0.02	9.67
CL	55.85	80.59	76.45	62.89	64.58	78.98	52.96	50.08
TRACIN	55.02	76.94	73.47	58.85	61.75	76.74	48.42	49.89
DEEP k -NN	56.21	82.35	75.24	63.08	57.40	70.69	56.75	63.85
SIMIFEAT-V	82.30	93.21	82.52	81.09	73.19	84.48	65.42	74.26
SIMIFEAT-R	83.28	95.56	83.58	82.26	74.67	88.68	62.89	73.53

Table 2. Comparisons of F_1 -scores (%). CORES, CL, TracIn: Use logit layers. SimiFeat-V/R: Use only representations. **All methods use the same fixed extractor from CLIP**. Top 2 are **bold**.

METHOD	CIFAR10				CIFAR100			
	Human	Symm. 0.6	Asym. 0.3	Inst. 0.4	Human	Symm. 0.6	Asym. 0.3	Inst. 0.4
CE SIEVE	67.21	94.56	5.24	8.41	16.24	88.55	2.6	1.63
CORES	83.18	96.94	12.05	88.89	38.52	92.33	7.02	85.52
CL	69.76	95.03	77.14	62.91	67.64	85.67	62.58	61.53
TRACIN	81.85	95.96	80.75	64.97	79.32	91.03	63.12	64.31
DEEP k -NN	82.98	87.47	76.96	77.42	72.33	82.95	64.96	74.25
SIMIFEAT-V	87.43	96.44	88.97	87.11	76.26	86.88	73.50	80.03
SIMIFEAT-R	87.45	96.74	89.04	91.14	79.21	90.54	68.14	77.37

confident learning (CL) (Northcutt et al., 2021a), TracIn (Pruthi et al., 2020), and deep k -NN (Bahri et al., 2020). We use ResNet34 (He et al., 2016) as the backbone network in this experiment.

Baseline settings All these three baselines require training a model with the noisy supervision. Specifically, CORES (Cheng et al., 2021a) trains ResNet34 on the noisy dataset and uses its proposed sample sieve to filter out the corrupted instances. We adopt its default setting during training and calculate the F_1 -score of the sieved out corrupted instances. Confident learning (CL) (Northcutt et al., 2021a) detects corrupted labels by firstly estimating probabilistic thresholds to characterize label noise, ranking instances based on model predictions, then filtering out corrupted instances based on ranking and thresholds. We adopt its default hyper-parameter setting to train ResNet34. TracIn (Pruthi et al., 2020) detects corrupted labels by evaluating the self-influence of each instance, where the corrupted instances tend to have a high influence score. The influence scores are calculated based on gradients of the last layer of ResNet34 at epoch 40, 50, 60, 100, where the model is trained with a batch size of 128. The initial learning rate is 0.1 and decays to 0.01 at epoch 50. Note TracIn only provides ranking for instances. To exactly detect corrupted instances, thresholds are required. For a fair comparison, we refer to the thresholds learned by confident learning (Northcutt et al., 2021a). Thus the corrupted instances selected by TracIn are based on the ranking from its self-influence and thresholds from CL. To highlight that our solutions work well without any supervision, our feature extractor $g(\cdot)$ comes from the ResNet34 pre-trained by SimCLR (Chen et al., 2020) where contrastive learning is applied and *no supervision* is

required. Extractor $g(\cdot)$ is obtained with only in-distribution features, e.g., for experiments with CIFAR-10, $g(\cdot)$ is pre-trained with features only from CIFAR-10. The detailed implementation for deep k -NN filter (Bahri et al., 2020) is not public. Noting their k -NN approach is employed on the logit layer, we reproduce their work by firstly training the model on the noisy data then substituting the model logits for $g(\cdot)$ in SimiFeat-V. The best epoch result for deep k -NN is reported.

Performance Table 1 compares the results obtained with or without supervisions. We can see both the voting-based and the ranking-based method achieve overall higher F_1 -scores compared with the other three results that require learning with noisy supervisions. Moreover, in detecting the real-world human-level noisy labels, our solution outperforms baselines around 20% on CIFAR-10 and 10% on CIFAR-100, which indicates the training-free solution are more robust to complicated noise patterns. One might also note that CORES achieves exceptionally low F_1 -scores on CIFAR-10/100 with asymmetric noise and CIFAR-100 with human noise. This observation also informs us that customized training processes might not be universally applicable.

5.2. Features May Be Better Than Model Predictions

Our next experiment aims to compare the performance of the data-centric method with the learning-centric method when the same feature extractor is adopted. Thus in this experiment, all methods adopt the same fixed feature extractor (ViT-B/32 pre-trained by CLIP (Radford et al., 2021)). Our proposed data-centric method directly operates on the extracted features, while the learning-centric method further

Table 3. Comparisons of F_1 -scores (%) using $g(\cdot)$ with different δ_k (%). Model names are the same as Figure 3.

PRE-TRAINED MODEL	CIFAR10			CIFAR100		
	$1 - \delta_k$	<i>Human</i>	<i>Inst. 0.4</i>	$1 - \delta_k$	<i>Human</i>	<i>Inst. 0.4</i>
R18-IMG	35.73	75.40	80.22	11.30	74.91	71.99
R34-IMG	48.13	79.52	82.43	16.17	76.88	74.00
R50-IMG	45.77	78.40	82.06	15.81	76.55	73.51
ViT-B/32-CLIP	64.12	87.45	91.14	19.94	79.21	77.37
R34-C10-SSL	69.31	83.28	85.26	2.59	68.03	65.94
R34-C10-CLEAN	99.41	98.39	98.59	0.22	60.90	60.73
R34-C100-SSL	18.59	59.96	74.99	22.46	74.67	73.53
R34-C100-CLEAN	18.58	60.17	76.41	89.07	92.87	95.29

 Table 4. Experiments on Clothing1M. None: Standard training with 1M noisy data. R50-Img (or ViT-B/32-CLIP, R50-Img Warmup-1): Apply our method with ResNet50 pre-trained on ImageNet (or ViT-B/32 pre-trained by CLIP, R50-Img with 1-epoch warmup). The clean test accuracy on the best epoch, the last 10 epochs, and the last epoch, are listed. Top-1 is **bold**.

DATA SELECTION	# TRAINING	BEST EPOCH	LAST 10	LAST
NONE	1M (100%)	70.32	69.44 \pm 0.13	69.53
R50-IMG	770K (77.0%)	72.37	71.95 \pm 0.08	71.89
ViT-B/32-CLIP	700K (70.0%)	72.54	72.23 \pm 0.17	72.11
R50-IMG WARMUP-1	767K (76.7%)	73.64	73.28 \pm 0.18	73.41

train a linear layer with noisy supervisions based on the extracted features. In addition to the baselines compared in Section 5.1, we also compare to CE Sieve (Cheng et al., 2021a) which follows the same sieving process as CORES but uses CE loss without regularizer. Other settings are the same as those in Section 5.1.

Table 2 summarizes the results of this experiment. By counting the frequency of reaching top-2 F_1 -scores, we find SimiFeat-R wins 1st place, SimiFeat-V and CORES are tied for 2nd place. However, similar to Table 2, we find the training process of CORES to be unstable. For instance, it almost fails for CIFAR-100 with asymmetric noise. Besides, comparing deep k -NN with SimiFeat-V, we find using the model logits given by an additional linear layer fine-tuned with noisy supervisions cannot always help improve the performance of detecting corrupted labels. It is therefore reasonable to believe both methods that directly deal with the extracted features achieve an overall higher F_1 -score than other learning-centric methods.

5.3. The Effect of the Quality of Features

Previous experiments demonstrate our methods overall outperform baselines with high-quality features. It is interesting to see how lower-quality features perform. We summarize results of SimiFeat-R in Table 3. There are several interesting findings: 1) The ImageNet pre-trained models perform well, indicating the traditional supervised training on out-of-distribution data helps obtained high-quality features; 2) For CIFAR-100, extractor $g(\cdot)$ obtained with only features

from CIFAR-10 (R34-C10-SSL) performs better than the extractor with clean CIFAR-10 (R34-C10-Clean), indicating that contrastive pre-training has better generalization ability to out-of-distribution data than supervised learning; 3) The F_1 -scores achieved by $g(\cdot)$ trained with the corresponding clean dataset are close to 1, indicating our solution can give perfect detection with ideal features.

5.4. More Experiments on Clothing1M

Besides, we test the performance of training only with the clean instances selected by our approach in Table 4. Standard training with Cross-Entropy loss is adopted. The only difference between the first row and other rows of Table 4 is that some training instances are filtered out by our approach. Table 4 shows simply filtering out corrupted instances based on our approach distinctively outperforms the baseline. We also observe that slightly tuning $g(\cdot)$ in the fine-grained Clothing1M dataset would be helpful. Note the best-epoch test accuracy we can achieve is 73.64%, which outperforms many baselines such as HOC 73.39% (Zhu et al., 2021b), GCE+SimCLR 73.35% (Ghosh & Lan, 2021), CORES 73.24% (Cheng et al., 2021a), GCE 69.75% (Zhang & Sabuncu, 2018). See more detailed settings and discussions in Appendix C.

6. Conclusions

This paper proposed a new and universally applicable data-centric training-free solution to detect noisy labels by using the neighborhood information of features. We have also demonstrated that the proposed data-centric method works even better than the learning-centric method when both methods are build on the same features. Future works will explore other tasks that could benefit from label cleaning, e.g., fairness (Liu & Wang, 2021) and multi-label learning (Liu et al., 2021c), and extend the idea to long-tail subpopulation detection (Liu, 2021; Wei et al., 2022a).

Acknowledgment This work is partially supported by the National Science Foundation (NSF) under grants IIS-2007951, IIS-2143895, and the Office of Naval Research under grant N00014-20-1-22.

References

- Agarwal, V., Podchiyska, T., Banda, J. M., Goel, V., Leung, T. I., Minty, E. P., Sweeney, T. E., Gyang, E., and Shah, N. H. Learning statistical models of phenotypes using noisy labeled training data. *Journal of the American Medical Informatics Association*, 23(6):1166–1173, 2016.
- Amid, E., Warmuth, M. K., Anil, R., and Koren, T. Robust bi-tempered logistic loss based on bregman divergences. In *Advances in Neural Information Processing Systems*, pp. 14987–14996, 2019.
- Bahri, D., Jiang, H., and Gupta, M. Deep k-nn for noisy labels. In *International Conference on Machine Learning*, pp. 540–550. PMLR, 2020.
- Bai, Y. and Liu, T. Me-momentum: Extracting hard confident examples from noisily labeled data. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 9312–9321, 2021.
- Bai, Y., Yang, E., Han, B., Yang, Y., Li, J., Mao, Y., Niu, G., and Liu, T. Understanding and improving early stopping for learning with noisy labels. *Advances in Neural Information Processing Systems*, 34, 2021.
- Berthelot, D., Carlini, N., Goodfellow, I., Papernot, N., Oliver, A., and Raffel, C. Mixmatch: A holistic approach to semi-supervised learning. *arXiv preprint arXiv:1905.02249*, 2019.
- Buldygin, V. V. and Kozachenko, Y. V. Sub-gaussian random variables. *Ukrainian Mathematical Journal*, 32(6): 483–489, 1980.
- Chen, T., Kornblith, S., Norouzi, M., and Hinton, G. A simple framework for contrastive learning of visual representations. In *International conference on machine learning*, pp. 1597–1607. PMLR, 2020.
- Cheng, H., Zhu, Z., Li, X., Gong, Y., Sun, X., and Liu, Y. Learning with instance-dependent label noise: A sample sieve approach. In *International Conference on Learning Representations*, 2021a. URL <https://openreview.net/forum?id=2VXyy9mIyU3>.
- Cheng, H., Zhu, Z., Sun, X., and Liu, Y. Demystifying how self-supervised features improve training from noisy labels. *arXiv preprint arXiv:2110.09022*, 2021b.
- Deng, J., Dong, W., Socher, R., Li, L.-J., Li, K., and Fei-Fei, L. ImageNet: A Large-Scale Hierarchical Image Database. In *CVPR09*, 2009.
- Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics*, pp. 4171–4186, June 2019. doi: 10.18653/v1/N19-1423. URL <https://aclanthology.org/N19-1423>.
- Dua, D. and Graff, C. UCI machine learning repository, 2017. URL <http://archive.ics.uci.edu/ml>.
- Feng, L., Shu, S., Lin, Z., Lv, F., Li, L., and An, B. Can cross entropy loss be robust to label noise? In *Proceedings of the Twenty-Ninth International Conference on International Joint Conferences on Artificial Intelligence*, pp. 2206–2212, 2021.
- Gao, W., Yang, B.-B., and Zhou, Z.-H. On the resistance of nearest neighbor to random noisy labels. *arXiv preprint arXiv:1607.07526*, 2016.
- Ghosh, A. and Lan, A. Contrastive learning improves model robustness under label noise. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 2703–2708, 2021.
- Ghosh, A., Kumar, H., and Sastry, P. Robust loss functions under label noise for deep neural networks. In *Thirty-First AAAI Conference on Artificial Intelligence*, 2017.
- Gong, M., Li, H., Meng, D., Miao, Q., and Liu, J. Decomposition-based evolutionary multiobjective optimization to self-paced learning. *IEEE Transactions on Evolutionary Computation*, 23(2):288–302, 2018.
- Han, B., Yao, Q., Yu, X., Niu, G., Xu, M., Hu, W., Tsang, I., and Sugiyama, M. Co-teaching: Robust training of deep neural networks with extremely noisy labels. In *Advances in neural information processing systems*, pp. 8527–8537, 2018.
- Han, B., Yao, Q., Liu, T., Niu, G., Tsang, I. W., Kwok, J. T., and Sugiyama, M. A survey of label-noise representation learning: Past, present and future. *arXiv preprint arXiv:2011.04406*, 2020.
- He, K., Zhang, X., Ren, S., and Sun, J. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770–778, 2016.
- He, K., Fan, H., Wu, Y., Xie, S., and Girshick, R. Momentum contrast for unsupervised visual representation learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 9729–9738, 2020.
- Huang, J., Qu, L., Jia, R., and Zhao, B. O2u-net: A simple noisy label detection approach for deep neural networks. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 3326–3334, 2019.

- Jaiswal, A., Babu, A. R., Zadeh, M. Z., Banerjee, D., and Makedon, F. A survey on contrastive self-supervised learning. *Technologies*, 9(1):2, 2021.
- Ji, X., Henriques, J. F., and Vedaldi, A. Invariant information clustering for unsupervised image classification and segmentation. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 9865–9874, 2019.
- Jiang, H., Kim, B., Guan, M., and Gupta, M. To trust or not to trust a classifier. *Advances in neural information processing systems*, 31, 2018.
- Jiang, L., Huang, D., Liu, M., and Yang, W. Beyond synthetic noise: Deep learning on controlled noisy labels. In *International Conference on Machine Learning*, pp. 4804–4815. PMLR, 2020.
- Jiang, Z., Zhou, K., Liu, Z., Li, L., Chen, R., Choi, S.-H., and Hu, X. An information fusion approach to learning with instance-dependent label noise. In *International Conference on Learning Representations*, 2022. URL <https://openreview.net/forum?id=ecH2FKaARUp>.
- Karger, D., Oh, S., and Shah, D. Iterative learning for reliable crowdsourcing systems. In *Neural Information Processing Systems*, NIPS ’11, 2011.
- Karger, D. R., Oh, S., and Shah, D. Efficient crowdsourcing for multi-class labeling. In *Proceedings of the ACM SIGMETRICS/international conference on Measurement and modeling of computer systems*, pp. 81–92, 2013.
- Kong, S., Li, Y., Wang, J., Rezaei, A., and Zhou, H. Knn-enhanced deep learning against noisy labels. *arXiv preprint arXiv:2012.04224*, 2020.
- Krizhevsky, A., Hinton, G., et al. Learning multiple layers of features from tiny images. Technical report, Citeseer, 2009.
- Krizhevsky, A., Sutskever, I., and Hinton, G. E. Imagenet classification with deep convolutional neural networks. *Advances in neural information processing systems*, 25: 1097–1105, 2012.
- Li, J., Socher, R., and Hoi, S. C. Dividemix: Learning with noisy labels as semi-supervised learning. In *International Conference on Learning Representations*, 2020a. URL <https://openreview.net/forum?id=HJgExaVtwr>.
- Li, J., Zhang, M., Xu, K., Dickerson, J., and Ba, J. How does a neural network’s architecture impact its robustness to noisy labels? *Advances in Neural Information Processing Systems*, 34, 2021a.
- Li, M., Soltanolkotabi, M., and Oymak, S. Gradient descent with early stopping is provably robust to label noise for overparameterized neural networks. In *International conference on artificial intelligence and statistics*, pp. 4313–4324. PMLR, 2020b.
- Li, X., Liu, T., Han, B., Niu, G., and Sugiyama, M. Provably end-to-end label-noise learning without anchor points. In *International Conference on Machine Learning*, pp. 6403–6413. PMLR, 2021b.
- Liu, A. T., Li, S.-W., and Lee, H.-y. Tera: Self-supervised learning of transformer encoder representation for speech. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 29:2351–2366, 2021a.
- Liu, E. Z., Haghighi, B., Chen, A. S., Raghunathan, A., Koh, P. W., Sagawa, S., Liang, P., and Finn, C. Just train twice: Improving group robustness without training group information. In *International Conference on Machine Learning*, pp. 6781–6792. PMLR, 2021b.
- Liu, Q., Peng, J., and Ihler, A. Variational inference for crowdsourcing. In *Proceedings of the 25th International Conference on Neural Information Processing Systems-Volume 1*, pp. 692–700, 2012.
- Liu, S., Niles-Weed, J., Razavian, N., and Fernandez-Granda, C. Early-learning regularization prevents memorization of noisy labels. *Advances in neural information processing systems*, 33:20331–20342, 2020.
- Liu, T. and Tao, D. Classification with noisy labels by importance reweighting. *IEEE Transactions on pattern analysis and machine intelligence*, 38(3):447–461, 2015.
- Liu, W., Wang, H., Shen, X., and Tsang, I. The emerging trends of multi-label learning. *IEEE transactions on pattern analysis and machine intelligence*, 2021c.
- Liu, Y. Understanding instance-level label noise: Disparate impacts and treatments. In *International Conference on Machine Learning*, pp. 6725–6735. PMLR, 2021.
- Liu, Y. Identifiability of label noise transition matrix. *arXiv preprint arXiv:2202.02016*, 2022.
- Liu, Y. and Guo, H. Peer loss functions: Learning from noisy labels without knowing noise rates. In *International Conference on Machine Learning*, pp. 6226–6236. PMLR, 2020.
- Liu, Y. and Liu, M. An online learning approach to improving the quality of crowd-sourcing. *ACM SIGMETRICS Performance Evaluation Review*, 43(1):217–230, 2015.
- Liu, Y. and Wang, J. Can less be more? when increasing-to-balancing label noise rates considered beneficial. *Advances in Neural Information Processing Systems*, 34, 2021.

- Luo, H., Cheng, H., Meng, F., Gao, Y., Li, K., Zhang, M., and Sun, X. An empirical study and analysis on open-set semi-supervised learning. *arXiv preprint arXiv:2101.08237*, 2021.
- Natarajan, N., Dhillon, I. S., Ravikumar, P. K., and Tewari, A. Learning with noisy labels. In *Advances in neural information processing systems*, pp. 1196–1204, 2013.
- Northcutt, C., Jiang, L., and Chuang, I. Confident learning: Estimating uncertainty in dataset labels. *Journal of Artificial Intelligence Research*, 70:1373–1411, 2021a.
- Northcutt, C. G., Athalye, A., and Mueller, J. Pervasive label errors in test sets destabilize machine learning benchmarks. In *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track (Round 1)*, 2021b. URL <https://openreview.net/forum?id=XccDXrDNLeK>.
- Patrini, G., Rozza, A., Krishna Menon, A., Nock, R., and Qu, L. Making deep neural networks robust to label noise: A loss correction approach. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 1944–1952, 2017.
- Pruthi, G., Liu, F., Kale, S., and Sundararajan, M. Estimating training data influence by tracing gradient descent. In *Advances in Neural Information Processing Systems*, volume 33, pp. 19920–19930, 2020.
- Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., et al. Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning*, pp. 8748–8763. PMLR, 2021.
- Reeve, H. and Kabán, A. Fast rates for a knn classifier robust to unknown asymmetric label noise. In *International Conference on Machine Learning*, pp. 5401–5409. PMLR, 2019.
- Shu, J., Zhao, Q., Chen, K., Xu, Z., and Meng, D. Learning adaptive loss for robust learning with noisy labels. *arXiv preprint arXiv:2002.06482*, 2020.
- Wang, H., Xiao, R., Li, Y., Feng, L., Niu, G., Chen, G., and Zhao, J. PiCO: Contrastive label disambiguation for partial label learning. In *International Conference on Learning Representations*, 2022a. URL <https://openreview.net/forum?id=EhYjZy6elgJ>.
- Wang, J., Guo, H., Zhu, Z., and Liu, Y. Policy learning using weak supervision. *Advances in Neural Information Processing Systems*, 34, 2021a.
- Wang, J., Liu, Y., and Levy, C. Fair classification with group-dependent label noise. In *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*, pp. 526–536, 2021b.
- Wang, J., Wang, X. E., and Liu, Y. Understanding instance-level impact of fairness constraints. In *International Conference on Machine Learning*. PMLR, 2022b.
- Wang, Y., Ma, X., Chen, Z., Luo, Y., Yi, J., and Bailey, J. Symmetric cross entropy for robust learning with noisy labels. In *Proceedings of the IEEE International Conference on Computer Vision*, pp. 322–330, 2019.
- Wang, Z., Jiang, J., Han, B., Feng, L., An, B., Niu, G., and Long, G. Seminll: A framework of noisy-label learning by semi-supervised learning. *arXiv preprint arXiv:2012.00925*, 2020.
- Wei, H., Feng, L., Chen, X., and An, B. Combating noisy labels by agreement: A joint training method with co-regularization. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13726–13735, 2020.
- Wei, H., Tao, L., Xie, R., and An, B. Open-set label noise can improve robustness against inherent label noise. *Advances in Neural Information Processing Systems*, 34, 2021a.
- Wei, H., Tao, L., Xie, R., Feng, L., and An, B. Open-sampling: Exploring out-of-distribution data for rebalancing long-tailed datasets. In *International Conference on Machine Learning (ICML)*. PMLR, 2022a.
- Wei, H., Xie, R., Cheng, H., Feng, L., An, B., and Li, Y. Mitigating neural network overconfidence with logit normalization. *arXiv preprint arXiv:2205.09310*, 2022b.
- Wei, H., Xie, R., Feng, L., Han, B., and An, B. Deep learning from multiple noisy annotators as a union. *IEEE Transactions on Neural Networks and Learning Systems*, 2022c.
- Wei, J. and Liu, Y. When optimizing f -divergence is robust with label noise. *arXiv preprint arXiv:2011.03687*, 2020.
- Wei, J., Liu, H., Liu, T., Niu, G., Sugiyama, M., and Liu, Y. To smooth or not? when label smoothing meets noisy labels, 2021b. URL <https://arxiv.org/abs/2106.04149>.
- Wei, J., Zhu, Z., Cheng, H., Liu, T., Niu, G., and Liu, Y. Learning with noisy labels revisited: A study using real-world human annotations. In *International Conference on Learning Representations*, 2022d. URL <https://openreview.net/forum?id=TBWA6PLJZQm>.
- Wei, J., Zhu, Z., Luo, T., Amid, E., Kumar, A., and Liu, Y. To aggregate or not? Learning with separate noisy

- labels, 2022e. URL <https://arxiv.org/abs/2206.07181>.
- Xia, X., Liu, T., Wang, N., Han, B., Gong, C., Niu, G., and Sugiyama, M. Are anchor points really indispensable in label-noise learning? In *Advances in Neural Information Processing Systems*, pp. 6838–6849, 2019.
- Xia, X., Liu, T., Han, B., Wang, N., Deng, J., Li, J., and Mao, Y. Extended T: Learning with mixed closed-set and open-set noisy labels. *arXiv preprint arXiv:2012.00932*, 2020a.
- Xia, X., Liu, T., Han, B., Wang, N., Gong, M., Liu, H., Niu, G., Tao, D., and Sugiyama, M. Part-dependent label noise: Towards instance-dependent label noise. In *Advances in Neural Information Processing Systems*, volume 33, pp. 7597–7610, 2020b.
- Xia, X., Liu, T., Han, B., Gong, C., Wang, N., Ge, Z., and Chang, Y. Robust early-learning: Hindering the memorization of noisy labels. In *International Conference on Learning Representations*, 2021.
- Xiao, T., Xia, T., Yang, Y., Huang, C., and Wang, X. Learning from massive noisy labeled data for image classification. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 2691–2699, 2015.
- Xie, Q., Dai, Z., Hovy, E., Luong, M.-T., and Le, Q. V. Unsupervised data augmentation. *arXiv preprint arXiv:1904.12848*, 2019.
- Yao, Q., Yang, H., Han, B., Niu, G., and Kwok, J. T. Searching to exploit memorization effect in learning with noisy labels. In *Proceedings of the 37th International Conference on Machine Learning, ICML ’20*, 2020.
- Yu, X., Han, B., Yao, J., Niu, G., Tsang, I., and Sugiyama, M. How does disagreement help generalization against label corruption? In *Proceedings of the 36th International Conference on Machine Learning*, volume 97, pp. 7164–7173. PMLR, 09–15 Jun 2019.
- Zhang, H., Cisse, M., Dauphin, Y. N., and Lopez-Paz, D. mixup: Beyond empirical risk minimization. In *International Conference on Learning Representations*, 2018. URL <https://openreview.net/forum?id=r1Ddp1-Rb>.
- Zhang, J., Sheng, V. S., Li, T., and Wu, X. Improving crowdsourced label quality using noise correction. *IEEE transactions on neural networks and learning systems*, 29(5):1675–1688, 2017.
- Zhang, M., Lee, J., and Agarwal, S. Learning from noisy labels with no change to the training process. In *International Conference on Machine Learning*, pp. 12468–12478. PMLR, 2021.
- Zhang, Y., Chen, X., Zhou, D., and Jordan, M. I. Spectral methods meet em: A provably optimal algorithm for crowdsourcing. *Advances in neural information processing systems*, 27:1260–1268, 2014.
- Zhang, Z. and Sabuncu, M. Generalized cross entropy loss for training deep neural networks with noisy labels. In *Advances in neural information processing systems*, pp. 8778–8788, 2018.
- Zhu, Z., Liu, T., and Liu, Y. A second-order approach to learning with instance-dependent label noise. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2021a.
- Zhu, Z., Song, Y., and Liu, Y. Clusterability as an alternative to anchor points when learning with noisy labels. In *Proceedings of the 38th International Conference on Machine Learning, ICML ’21*, 2021b.
- Zhu, Z., Zhu, J., Liu, J., and Liu, Y. Federated bandit: A gossiping approach. In *Abstract Proceedings of the 2021 ACM SIGMETRICS/International Conference on Measurement and Modeling of Computer Systems*, pp. 3–4, 2021c.
- Zhu, Z., Luo, T., and Liu, Y. The rich get richer: Disparate impact of semi-supervised learning. In *International Conference on Learning Representations*, 2022a. URL <https://openreview.net/forum?id=DXPftn5kjQK>.
- Zhu, Z., Wang, J., and Liu, Y. Beyond images: Label noise transition matrix estimation for tasks with lower-quality features. *arXiv preprint arXiv:2202.01273*, 2022b.

The omitted proofs and experiment settings are provided as follows.

A. Theoretical Analyses

A.1. Proof for Proposition 4.1

Now we derive a lower bound for the probability of getting true detection with majority vote:

$$\begin{aligned} \mathbb{P}(\text{Vote is correct}|k) &\geq (1 - \delta_k) \cdot \left[p \sum_{l=0}^{\lceil (k+1)/2 \rceil - 1} \binom{k+1}{l} e^l (1-e)^{k+1-l} \right. \\ &\quad \left. + (1-p) \sum_{l=0}^{\lceil (k+1)/2 \rceil - 1} \binom{k+1}{l} e^l (1-e)^{k+1-l} \right] \\ &= (1 - \delta_k) \cdot [p \cdot I_{1-e}(k+1-k', k'+1) + (1-p) \cdot I_{1-e}(k+1-k', k'+1)] \end{aligned}$$

where $I_{1-e}(k+1-k', k'+1)$ is the regularized incomplete beta function defined as

$$I_{1-e}(k+1-k', k'+1) = (k+1-k') \binom{k+1}{k'} \int_0^{1-e} t^{k-k'} (1-t)^{k'} dt,$$

and $k' = \lceil (k+1)/2 \rceil - 1$.

B. Proof for Theorem 4.3

Proof. Now we derive the worst-case error bound. We first repeat the notations defined in Section 4.2 as follows.

Denote random variable S by the score of each instance being clean. A higher score S indicates the instance is more likely to be clean. Denote the score of a true/false instance by

$$\begin{aligned} S_{n,j}^{\text{true}} &:= \text{Score}(\hat{y}_n, j) \mid (\tilde{y}_n = y_n = j), \\ S_{n',j}^{\text{false}} &:= \text{Score}(\hat{y}_{n'}, j) \mid (\tilde{y}_{n'} = j, y_{n'} \neq j). \end{aligned}$$

Both are scalars. Then for instances in \mathcal{N}_j , we have two set of random variables $\mathbb{S}_j^{\text{true}} := \{S_{n,j}^{\text{true}} \mid n \in \mathcal{N}_j, \tilde{y}_n = y_n = j\}$ and $\mathbb{S}_j^{\text{false}} := \{S_{n',j}^{\text{false}} \mid n' \in \mathcal{N}_j, \tilde{y}_{n'} = j, y_{n'} \neq j\}$. Recall $\mathcal{N}_j := \{n \mid \tilde{y}_n = j\}$ are the set of indices that correspond to noisy class j . Intuitively, the score $S_{n,j}^{\text{true}}$ should be greater than $S_{n',j}^{\text{false}}$. Suppose their means, which depend on noise rates, are bounded, i.e.,

$$\mathbb{E}[S_{n,j}^{\text{true}}] \geq \mu_j^{\text{true}}, \quad \mathbb{E}[S_{n',j}^{\text{false}}] \leq \mu_j^{\text{false}}$$

for all feasible n, n' . Assume there exists a feasible v such that both S_j^{true} and S_j^{false} follow sub-Gaussian distributions with variance proxy $\frac{\Delta^2}{2v}$ (Buldygin & Kozachenko, 1980; Zhu et al., 2021c) such that:

$$\mathbb{P}(\mu_j^{\text{true}} - S_{n,j}^{\text{true}} \geq t) \leq e^{-\frac{vt^2}{\Delta^2}}, \quad \mathbb{P}(S_{n',j}^{\text{false}} - \mu_j^{\text{false}} \geq t) \leq e^{-\frac{vt^2}{\Delta^2}},$$

and the probability density satisfies $\mathbb{P}(S_j^{\text{true}} = \mu_j^{\text{true}}) = \mathbb{P}(S_j^{\text{false}} = \mu_j^{\text{false}}) = 1/\Delta$, where $1/\Delta$ is the ‘‘height’’ of both distributions, v is the decay rate of tails. Let N_j^- (N_j^+) be the number of indices in $\mathbb{S}_j^{\text{false}}$ ($\mathbb{S}_j^{\text{true}}$).

For ease of notations, we omit the subscript j in this proof since the detection is performed on each j individually.

Let S^{false} (S^{true}) be an arbitrary random variable in $\mathbb{S}^{\text{false}}$ (\mathbb{S}^{true}). Denote the order statistics of random variables in set $\mathbb{S}^{\text{false}}$ by $S_{(1)}^{\text{false}}, \dots, S_{(N^-)}^{\text{false}}$, where $S_{(1)}^{\text{false}}$ is the smallest order statistic and $S_{(N^-)}^{\text{false}}$ is the largest order statistic. The following lemma motivates the performance of the rank-based method.

Lemma B.1. *The F_1 -score of detecting corrupted labels in \mathcal{N}_j by the rank-based method will be no less than $1 - \alpha/N^-$ when the true probability $\mathbb{P}(Y = j \mid \tilde{Y} = j)$ is known and $S_{(N^-)}^{\text{false}} < S_{(\alpha+1)}^{\text{true}}$.*

Lemma B.1 connects the upper bound for the number of wrongly detected corrupted instances with order statistics. There are two cases that can cause detection errors:

Case-1:

$$0 \leq \mu^{\text{true}} - S^{\text{true}} < \Delta \text{ and } 0 \leq S^{\text{false}} - \mu^{\text{false}} < \Delta : \text{at most } \alpha \text{ errors when } S_{(N^-)}^{\text{false}} < S_{(\alpha+1)}^{\text{true}}.$$

and **Case-2:**

$$\mu^{\text{true}} - S^{\text{true}} \geq \Delta \text{ or } S^{\text{false}} - \mu^{\text{false}} \geq \Delta : \text{at most } \max(N^-, N^+) \text{ errors}$$

We analyze each case as follows.

Case-1: When Case-1 holds, we have

$$\mathbb{P}(\mu^{\text{true}} - S^{\text{true}} = x) \leq 1/\Delta, x \in [0, \Delta]$$

and

$$\mathbb{P}(S^{\text{false}} - \mu^{\text{true}} = x) \leq 1/\Delta, x \in [0, \Delta].$$

The above two inequalities show that the left tail of S^{true} and the right tail of S^{false} can be upper bounded by uniform distributions. Denote the corresponding uniform distribution by $U^{\text{true}} \sim \text{Unif}(\mu^{\text{true}} - \Delta, \mu^{\text{true}})$ and $U^{\text{false}} \sim \text{Unif}(\mu^{\text{false}}, \mu^{\text{false}} + \Delta)$.

With true $\mathbb{P}(Y = j | \tilde{Y} = j)$, the detection errors only exist in the cases when the left tail of S^{true} and the right tail of S^{false} are overlapped. When the tails are upper bounded by uniform distributions, we have

$$\begin{aligned} \mathbb{P}(S_{(N^-)}^{\text{false}} < S_{(\alpha+1)}^{\text{true}}) &\geq \mathbb{P}(U_{(N^-)}^{\text{false}} < U_{(\alpha+1)}^{\text{true}}) \\ &= \mathbb{P}\left([U^{\text{false}} - \mu^{\text{false}}]_{(N^-)} + \mu^{\text{false}} < [U^{\text{true}} - (\mu^{\text{true}} - \Delta)]_{(\alpha+1)} + (\mu^{\text{true}} - \Delta)\right) \\ &= \mathbb{P}\left([U^{\text{false}} - \mu^{\text{false}}]_{(N^-)} - [U^{\text{true}} - (\mu^{\text{true}} - \Delta)]_{(\alpha+1)} < \mu^{\text{true}} - \mu^{\text{false}} - \Delta\right). \end{aligned}$$

Note

$$[U^{\text{false}} - \mu^{\text{false}}]_{(N^-)} \sim \text{Beta}(N^-, 1),$$

and

$$[U^{\text{true}} - (\mu^{\text{true}} - \Delta)]_{(\alpha+1)} \sim \text{Beta}(\alpha + 1, N^+ - \alpha),$$

where Beta denotes the Beta distribution. Both variables are independent. Thus the PDF of the difference is

$$f(p) = \begin{cases} B(N^+ - \alpha, 1)p^{N^+ - \alpha}(1-p)^{\alpha+1}F(1, N^- + N^+, 1 - N^-; \alpha + 2; 1 - p, 1 - p^2)/A, & 0 < p \leq 1 \\ B(N^- - \alpha, N^+ - \alpha)(-p)^{N^+ - \alpha}(1+p)^{N^- + N^+ - \alpha - 1}F(N^+ - \alpha, -\alpha, N^- + N^+; N^- + N^+ - \alpha; 1 - p^2, 1 + p)/A, & -1 \leq p < 0 \\ p = 0, & p = 0, \end{cases}$$

where $A = B(N^-, 1)B(\alpha + 1, N^+ - \alpha)$, $B(a, b) = \int_0^1 t^{a-1}(1-t)^{b-1}dt$

$$F(a, b_1, b_2; c; x, y) = \frac{\Gamma(c)}{\Gamma(a)\Gamma(c-a)} \int_0^1 t^{a-1}(1-t)^{c-a-1}(1-xt)^{-b_1}(1-yt)^{-b_2} dt.$$

Therefore, we have

$$\mathbb{P}(S_{(N^-)}^{\text{false}} < S_{(\alpha+1)}^{\text{true}}) \geq \int_{-1}^{\mu^{\text{true}} - \mu^{\text{false}} - \Delta} f(p) dp.$$

Case-2 The other part, we have no more than $e^{-v} \cdot \max(N^-, N^+)$ corrupted instances that may have higher scores than one clean instance.

Wrap-up From the above analyses, we know, w.p. at least $\int_{-1}^{\mu^{\text{true}} - \mu^{\text{false}} - \Delta} f(p) dp$, there are at most $e^{-v} \max(N^-, N^+) + \alpha$ errors in detection corrupted instances. Note Precision = Recall if we detect with the best threshold $N_j \mathbb{P}(Y = j | \tilde{Y} = j)$. Therefore, the corresponding F_1 -score would be at least $1 - \frac{e^{-v} \max(N^-, N^+) + \alpha}{N^-}$.

□

Table 5. Experiments on Clothing1M (Xiao et al., 2015) with or without balanced sampling. None: Standard training with 1M noisy data. R50-Img (or ViT-B/32-CLIP, R50-Img Warmup-1): Apply our method with ResNet50 pre-trained on ImageNet (or ViT-B/32 pre-trained by CLIP, R50-Img with 1-epoch warmup).

DATA SELECTION	# TRAINING SAMPLES	BEST EPOCH	LAST 10 EPOCHS	LAST EPOCH
NONE (STANDARD BASELINE) (UNBALANCED)	1M (100%)	70.32	69.44 \pm 0.13	69.53
NONE (STANDARD BASELINE) (BALANCED)	1M (100%)	72.20	71.40 \pm 0.31	71.22
R50-IMG (UNBALANCED)	770K (77.0%)	72.37	71.95 \pm 0.08	71.89
R50-IMG (BALANCED)	770K (77.0%)	72.42	72.06 \pm 0.16	72.24
ViT-B/32-CLIP (UNBALANCED)	700K (70.0%)	72.54	72.23 \pm 0.17	72.11
ViT-B/32-CLIP (BALANCED)	700K (70.0%)	72.99	72.76 \pm 0.15	72.91
R50-IMG WARMUP-1 (UNBALANCED)	767K (76.7%)	73.64	73.28 \pm 0.18	73.41
R50-IMG WARMUP-1 (BALANCED)	767K (76.7%)	73.97	73.37 \pm 0.03	73.35

C. Experiment Settings on Clothing1M

We firstly perform noise detection on 1 million noisy training instances then train only with the selected clean data to check the effectiveness. Particularly, in each epoch of the noisy detection, we use a batch size of 32 and sample 1,000 mini-batches from 1M training instances while ensuring the (noisy) labels are balanced. We repeat noisy detection for 600 epochs to ensure a full coverage of 1 million training instances. Parameter k is set to 10.

Feature Extractor: We tested three different feature extractors in Table 4: R50-Img, ViT-B/32-CLIP, and R50-Img Warmup-1. The former two feature extractors are the same as the ones used in Table 3. Particularly, R50-Img means the feature extractor is the standard ResNet50 encoder (removing the last linear layer) pre-trained on ImageNet (Deng et al., 2009). ViT-B/32-CLIP indicates the feature extractor is a vision transformer pre-trained by CLIP (Radford et al., 2021). Noting that Clothing1M is a fine-grained dataset. To get better domain-specific fine-grained visual features, we slightly train the ResNet50 pre-trained with ImageNet for one epoch, i.e., 1,000 mini-batches (batch size 32) randomly sampled from 1M training instances while ensuring the (noisy) labels are balanced. The learning rate is 0.002.

Training with the selected clean instances: Given the selected clean instances from our approach, we directly apply the Cross-Entropy loss to train a ResNet50 initialized by standard ImageNet pre-trained parameters. We **did not** apply any sophisticated training techniques, e.g., mixup (Zhang et al., 2018), dual networks (Li et al., 2020a; Han et al., 2018), loss-correction (Liu & Tao, 2015; Natarajan et al., 2013; Patrini et al., 2017), and robust loss functions (Liu & Guo, 2020; Cheng et al., 2021a; Zhu et al., 2021a; Wei & Liu, 2020). We train the model for 80 epochs with a batch size of 32. We sample 1,000 mini-batches per epoch randomly selected from 1M training instances. Note Table 4 does not apply balanced sampling. Only the pure cross-entropy loss is applied. We also test the performance with balanced training, i.e., in each epoch, ensure the noisy labels from each class are balanced. Our approach can be consistently benefited by balanced training, and achieves an accuracy of 73.97 in the best epoch, outperforming many baselines such as HOC 73.39% (Zhu et al., 2021b), GCE+SimCLR 73.35% (Ghosh & Lan, 2021), CORES 73.24% (Cheng et al., 2021a), GCE 69.75% (Zhang & Sabuncu, 2018). We believe the performance could be further improved by using some sophisticated training techniques mentioned above.